ASSIGNMENT-2

Assignment Date	3rd October 2022
Name	Muralidharan.M
Rollnumber	820319205022
Maximum Marks	2Marks

```python
import pandas as pd

import seaborn as sns

import matplotlib.pyplot as plt

import numpy as np

sns.set_style('darkgrid')

sns.set(font_scale=1.3)
```

In [25]:

```python
df=pd.read_csv("/content/drive/MyDrive/IBM/Assignment - 2 /Churn_Modelling.csv")
```

In [26]:

```python
df.head ()
```

Out [26]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

In [29]:

```python
df.drop (["RowNumber","CustomerId","Surname"], axis=1, inplace=True)
```

In [30]:

```python
df.info ()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
 #  Column         Non-Null Count  Dtype
--- ------         --------------  -----
 0  CreditScore    10000 non-null  int64
 1  Geography      10000 non-null  object
 2  Gender         10000 non-null  object
 3  Age            10000 non-null  int64
 4  Tenure         10000 non-null  int64
 5  Balance        10000 non-null  float64
 6  NumOfProducts  10000 non-null  int64
 7  HasCrCard      10000 non-null  int64
```

8   IsActiveMember   10000 non-null int64

 9   EstimatedSalary 10000 non-null float64

 10 Exited   10000 non-null int64

dtypes: float64(2), int64(7), object(2)

memory usage: 859.5+ KB

In [28]:

#Perform Univariate Analysis

plt**.**figure(figsize**=**(8,8))

sns**.**countplot(x**=**'Tenure',data**=**df)

plt**.**xlabel('0:Customers with Bank, 1: exited from bank')

plt**.**ylabel('No.of.Customers')
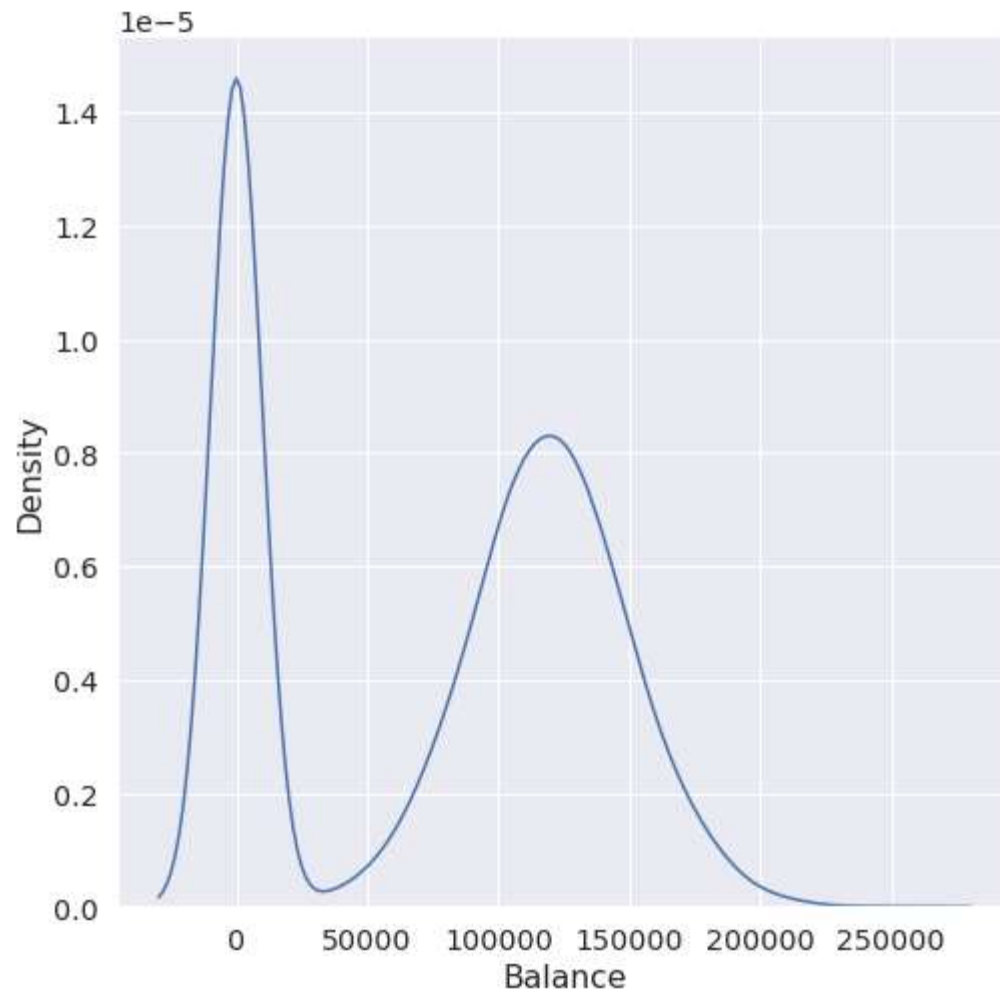
plt**.**title("Bank Customers viz")

plt**.**show()

Bank Customers viz
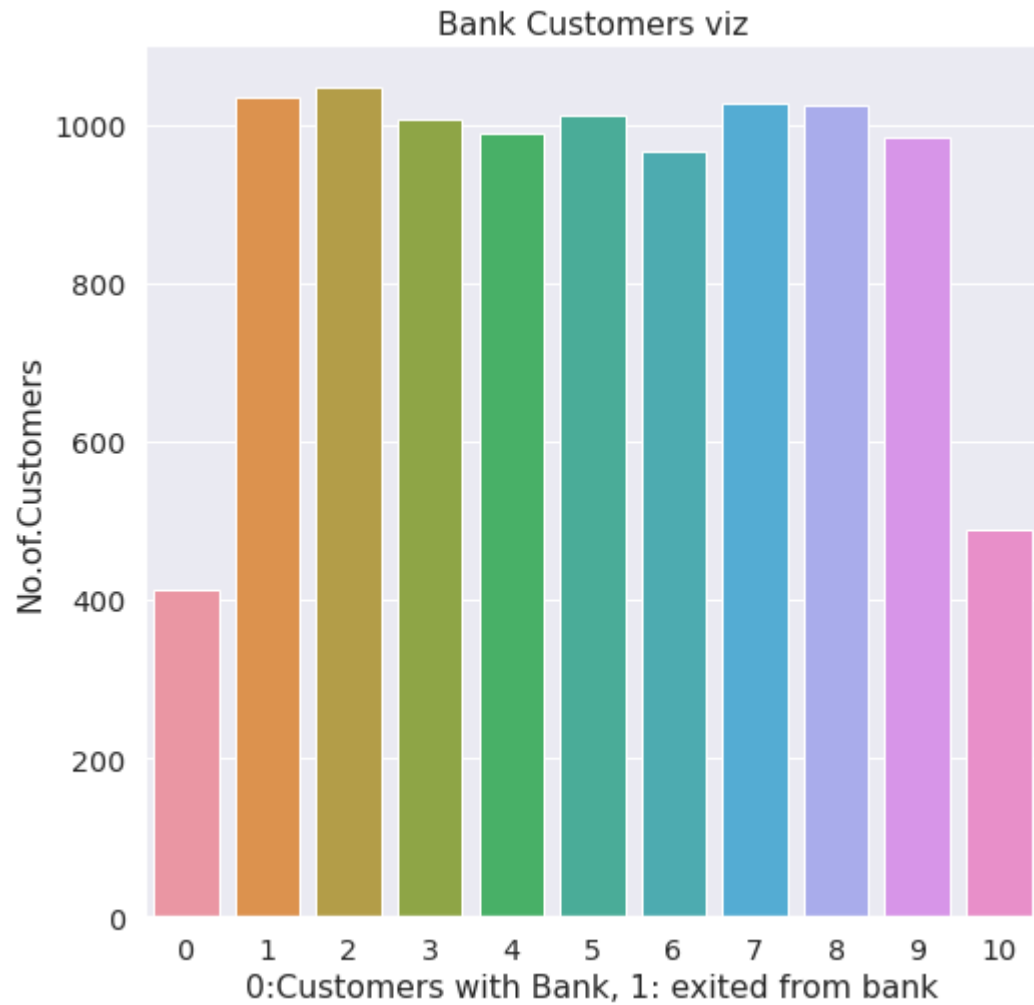
In[9]#Perform Univariate Analysis

plt.figure(figsize=(8,8))

sns.kdeplot(x=df['Balance'])

Out[9]:

<matplotlib.axes._subplots.AxesSubplot at 0x7fa0c03906d0>

In[10]*#Perform Bivariate Analysis*

plt.scatter(df.Age,df.Balance)

Out[10]:

<matplotlib.collections.PathCollection at 0x7fa0d35a7dd0>
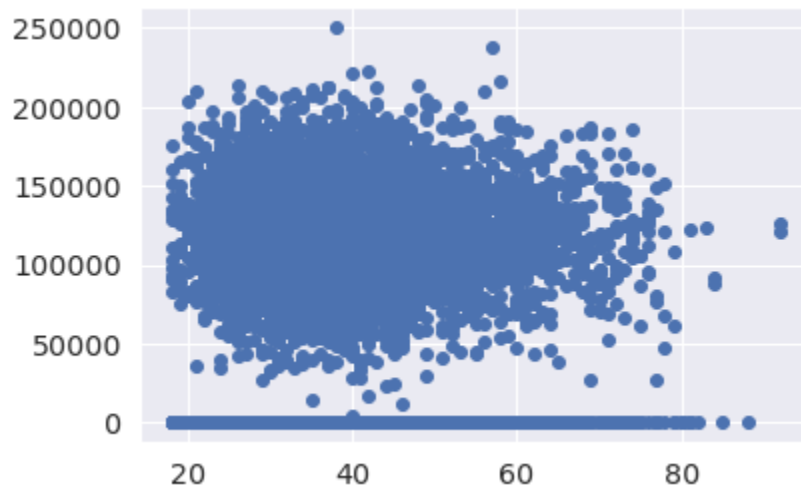
#Perform Bivariate Analysis

df.corr()

Out[54]:

	Credit Score	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActive Member	EstimatedSalary	Exited
CreditScore	1.000000	0.007888	-0.003965	0.000842	0.006268	0.012238	-0.005458	0.025651	-0.001384	-0.027094
Gender	0.007888	1.000000	0.022812	0.003739	0.069408	0.003972	-0.008523	0.006724	-0.001369	0.035943
Age	-0.003965	0.022812	1.000000	-0.009997	0.028308	-0.030680	-0.011721	0.085472	-0.007201	0.285323
Tenure	0.000842	0.003739	-0.009997	1.000000	-0.012254	0.013444	0.022583	-0.028362	0.007784	-0.014001
Balance	0.006268	0.069408	0.028308	-0.012254	1.000000	-0.304180	-0.014858	-0.010084	0.012797	0.118533

	Credit Score	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
NumOfProducts	0.012238	0.003972	-0.030680	0.013444	-0.304180	1.000000	0.003183	0.009612	0.014204	-0.047820
HasCrCard	-0.005458	-0.008523	-0.011721	0.022583	-0.014858	0.003183	1.000000	-0.011866	-0.009933	-0.007138
IsActiveMember	0.025651	0.006724	0.085472	-0.028362	-0.010084	0.009612	-0.011866	1.000000	-0.011421	-0.156128
EstimatedSalary	-0.001384	-0.001369	-0.007201	0.007784	0.012797	0.014204	-0.009933	-0.011421	1.000000	0.012097
Exited	-0.027094	0.035943	0.285323	-0.014001	0.118533	-0.047820	-0.007138	-0.156128	0.012097	1.000000

In [36]:

#Perform Bivariate Analysis

import statsmodels.api **as** sm

#define response variable

y = df['CreditScore']

#define explanatory variable

x = df[['EstimatedSalary']]

#add constant to predictor variables

x = sm**.**add_constant(x)

#fit linear regression model

model = sm.OLS(y, x).fit()

#view model summary

print(model.summary())

OLS Regression Results

===

Dep. Variable:	CreditScore	R-squared:	0.000
Model:	OLS	Adj. R-squared:	-0.000
Method:	Least Squares	F-statistic:	0.01916
Date:	Sat, 24 Sep 2022	Prob (F-statistic):	0.890
Time:	05:06:19	Log-Likelihood:	-59900.
No. Observations:	10000	AIC:	1.198e+05
Df Residuals:	9998	BIC:	1.198e+05
Df Model:	1		
Covariance Type:	nonrobust		

==

	coef	std err	t	P>|t|	[0.025	0.975]
const	650.7617	1.940	335.407	0.000	646.958	654.565
EstimatedSalary	-2.326e-06	1.68e-05	-0.138	0.890	-3.53e-05	3.06e-05

===

Omnibus:	132.939	Durbin-Watson:	2.014
Prob(Omnibus):	0.000	Jarque-Bera (JB):	84.242
Skew:	-0.072	Prob(JB):	5.10e-19
Kurtosis:	2.574	Cond. No.	2.32e+05

===

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.32e+05. This might indicate that there are

strong multicollinearity or other numerical problems.

/usr/local/lib/python3.7/dist-packages/statsmodels/tsa/tsatools.py:142: FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only

```
  x = pd.concat(x[::order], 1)
```

In [35]:

#Perform Multivariate Analysis

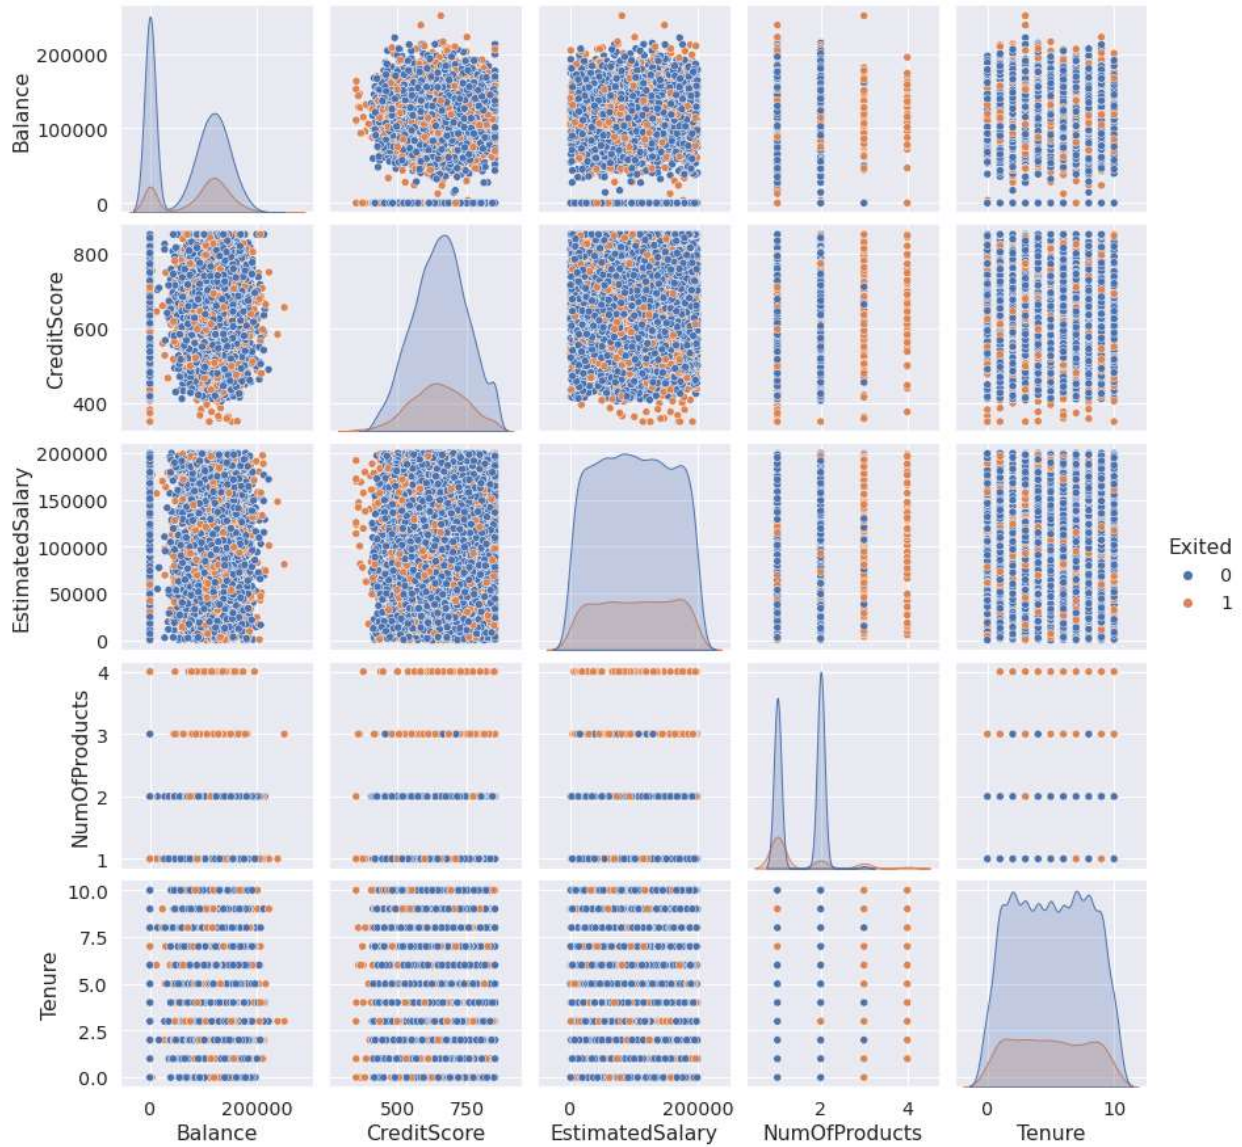plt**.**figure(figsize**=**(4,4))

sns**.**pairplot(data=df[["Balance","CreditScore","EstimatedSalary","NumOfProducts","Tenure","Exited"]], hue="Exited")

Out[35]:

<seaborn.axisgrid.PairGrid at 0x7fa0b00a1b10>

<Figure size 288x288 with 0 Axes>

#Perform Descriptive Statistics

df=pd.DataFrame(df)

print(df.sum())

CreditScore   6505288

Geography   FranceSpainFranceFranceSpainSpainFranceGermany...

Gender   FemaleFemaleFemaleFemaleFemaleMaleMaleFemaleMa...

Age   389218

Tenure   50128

Balance   764858892.88

NumOfProducts	15302
HasCrCard	7055
IsActiveMember	5151
EstimatedSalary	1000902398.81
Exited	2037

dtype: object

In [39]:

#Perform Descriptive Statistics

```python
print("----Sum Value-----")
print(df.sum(1))
print("-------------------------------")
print("-----Product Value-----")
print(df.prod())
print("-------------------------------")
```

----Sum Value-----

0	102015.88
1	197002.44
2	274149.37
3	94567.63
4	205492.92
	...
9995	97088.64
9996	159633.38
9997	42840.58
9998	168784.83
9999	169159.57

Length: 10000, dtype: float64

-----Product Value-----

CreditScore   0.0

Age   0.0

Tenure   0.0

Balance   0.0

NumOfProducts   0.0

HasCrCard   0.0

IsActiveMember   0.0

EstimatedSalary   inf

Exited   0.0

dtype: float64

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:3: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

  This is separate from the ipykernel package so we can avoid doing imports until

/usr/local/lib/python3.7/dist-packages/numpy/core/_methods.py:52: RuntimeWarning: overflow encountered in reduce

  return umr_prod(a, axis, dtype, out, keepdims, initial, where)

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:6: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

In [38]:

#Perform Descriptive Statistics

print("----------Mean Value-----------")

print(df.mean())

print("------------------------------")

print("----------Median Value---------")

print(df.median())

print("-----------------------------")

```
print("----------Mode Value------------")
```

```
print(df.mode())
```

```
print("------------------------------")
```

----------Mean Value-----------

CreditScore   650.528800

Age   38.921800

Tenure   5.012800

Balance   76485.889288

NumOfProducts   1.530200

HasCrCard   0.705500

IsActiveMember   0.515100

EstimatedSalary   100090.239881

Exited   0.203700

dtype: float64

----------Median Value---------

CreditScore   652.000

Age   37.000

Tenure   5.000

Balance   97198.540

NumOfProducts   1.000

HasCrCard   1.000

IsActiveMember   1.000

EstimatedSalary   100193.915

Exited   0.000

dtype: float64

----------Mode Value------------

   CreditScore Geography Gender Age Tenure Balance NumOfProducts \

```
0    850   France  Male  37    2    0.0         1
```

HasCrCard IsActiveMember EstimatedSalary Exited

```
0    1         1         24924.92     0
```

In [41]:

#Handling with missing Values

df**.**isnull()*#Checking values are null*

Out[41]:

	Credit Score	**Geogr aphy**	**Gen der**	**Ag e**	**Ten ure**	**Bala nce**	**NumOfPr oducts**	**HasCr Card**	**IsActiveM ember**	**Estimate dSalary**	**Exit ed**
0	False	False	Fals e	Fal se	Fals e	Fals e	False	False	False	False	Fals e
1	False	False	Fals e	Fal se	Fals e	Fals e	False	False	False	False	Fals e
2	False	False	Fals e	Fal se	Fals e	Fals e	False	False	False	False	Fals e
3	False	False	Fals e	Fal se	Fals e	Fals e	False	False	False	False	Fals e
4	False	False	Fals e	Fal se	Fals e	Fals e	False	False	False	False	Fals e

	Credit Score	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
...	...	...	...	...	...	...	...	...	...	...	...
9995	False	False	False	False	False	False	False	False	False	False	False
9996	False	False	False	False	False	False	False	False	False	False	False
9997	False	False	False	False	False	False	False	False	False	False	False
9998	False	False	False	False	False	False	False	False	False	False	False
9999	False	False	False	False	False	False	False	False	False	False	False

10000 rows × 11 columns

In [42]:

#Handling with missing Values

df**.**notnull()*#Checking values are not null*

Out[42]:

	Credit Score	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	True	True	True	True	True	True	True	True	True	True	True
1	True	True	True	True	True	True	True	True	True	True	True
2	True	True	True	True	True	True	True	True	True	True	True

	Credit Score	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
3	True	True	True	True	True	True	True	True	True	True	True
4	True	True	True	True	True	True	True	True	True	True	True
...	...	...	...	...	...	...	...	...	...	...	...
9995	True	True	True	True	True	True	True	True	True	True	True
9996	True	True	True	True	True	True	True	True	True	True	True
9997	True	True	True	True	True	True	True	True	True	True	True
9998	True	True	True	True	True	True	True	True	True	True	True
9999	True	True	True	True	True	True	True	True	True	True	True

10000 rows × 11 columns

In [43]:

#Find outliers & replace the outliers

sns**.**boxplot(df['Balance'])

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

  FutureWarning

Out[43]:

<matplotlib.axes._subplots.AxesSubplot at 0x7fa0af6dcf90>

#Find outliers & replace the outliers

print(np**.**where(df['Balance']**>**100000))

(array([2,   4,   5, ..., 9987, 9993, 9999]),)

9995	771	France	0	39	5	0.00	2	1	0	96270.64	0
9996	516	France	0	35	10	57369.61	1	1	1	101699.77	0
9997	709	France	0	36	7	0.00	1	0	1	42085.58	1
9998	772	Germany	1	42	3	75075.31	2	1	0	92888.52	1
9999	792	France	0	28	4	130142.79	1	1	0	38190.78	0

10000 rows × 11 columns

3	39	1	0.00
4	43	2	125510.82
...	...	...	...
9995	39	5	0.00
9996	35	10	57369.61
9997	36	7	0.00
9998	42	3	75075.31
9999	28	4	130142.79

[10000 rows x 3 columns]

---------Independent Variables---------

0	1
1	1
2	3
3	2
4	1

..

9995   2

9996   1

9997   1

9998   2

9999   1

Name: NumOfProducts, Length: 10000, dtype: int64

In [ ]:

#Scale the independent Variables

from sklearn.preprocessing **import** StandardScaler

object**=** StandardScaler()

 # standardization

scale **=** object**.**fit_transform(df)

print(scale)

[[-0.32622142 0.29351742 -1.04175968 ... 0.97024255 0.02188649

   1.97716468]

 [-0.44003595 0.19816383 -1.38753759 ... 0.97024255 0.21653375

  -0.50577476]

 [-1.53679418 0.29351742 1.03290776 ... -1.03067011 0.2406869

   1.97716468]

 ...

 [0.60498839 -0.27860412 0.68712986 ... 0.97024255 -1.00864308

   1.97716468]

 [1.25683526 0.29351742 -0.69598177 ... -1.03067011 -0.12523071

   1.97716468]

 [1.46377078 -1.04143285 -0.35020386 ... -1.03067011 -1.07636976

  -0.50577476]]

In [ ]:

#Split the data into training & testing

from sklearn.model_selection **import** train_test_split

In [ ]:

#Split the data into training & testing

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=4,random_state=4)

x_train

Out[ ]:

	const	EstimatedSalary
2558	1.0	137903.54
7642	1.0	121765.00
8912	1.0	109470.34
3319	1.0	2923.61
6852	1.0	7312.25
...	...	...
456	1.0	7666.73
6017	1.0	9085.00
709	1.0	147794.63
8366	1.0	102515.42
1146	1.0	54776.64

9996 rows × 2 columns

In [ ]:

#Split the data into training & testing

x_test

Out[ ]:

	const	EstimatedSalary
1603	1.0	23305.85
8713	1.0	41248.80
4561	1.0	143317.42
6600	1.0	174123.16

In [ ]:

#Split the data into training & testing

y_train

Out[ ]:

2558   727

7642   811

8912   623

3319   430

6852   600

   ...

456   733

6017   487

709   686

8366   637

1146   614

Name: CreditScore, Length: 9996, dtype: int64

In [ ]:

#Split the data into training & testing

y_test

Out[ ]:

1603   576

8713   786

4561   562

6600   505

Name: CreditScore, dtype: int64